



KubeCon



CloudNativeCon

China 2018

# Stop hitting yourself!

Michael Russell





KubeCon



CloudNativeCon

China 2018

“ ”

**Computers do what you **say**,**  
**Software should do what you **mean**.**

- Me



KubeCon



CloudNativeCon

China 2018

“ ”

**Software is easy to use**  
**but hard to run.**

- Also me



# About me



KubeCon



CloudNativeCon

China 2018

Australian living in the Netherlands

Works at Elastic

Responds to: Michael, Mick, Micky, Mike,  
Mikey, Crazybus, Rusty

Likes: Food, travelling, gaming



# Stopping hitting yourself!



KubeCon



CloudNativeCon

China 2018



# Disclaimer



KubeCon



CloudNativeCon

China 2018



A lot of this talk is going to come off pretty negative

I actually do like Kubernetes

I didn't get paid to say that

I didn't get paid at all actually

# Restore



KubeCon



CloudNativeCon

China 2018

**Dev**

"Are you busy?" (yes)

**Dev**

"We need to restore the production database." (Oh noes)

**Me**

"Do we have backups?"

**Dev**

"Yeah there is a bash script and a cronjob somewhere."



# Bash bashing



KubeCon



CloudNativeCon

China 2018

```
#!/bin/bash
# Author: Some dude
#
# Changelog
# 1995/05/23 - Some dude - Added mysql backups for important company database

mysqldump --all-databases > /mnt/backup/$(date +%s).sql
```



# 20 years later...



KubeCon



CloudNativeCon

China 2018

```
#!/bin/bash
# Author: Some dude
#
# Changelog
# 1995/05/23 - Some dude - Added mysql backups for important company database
# 1995/05/27 - Some dude - Added debug logging and log backup duration
# 1998/02/13 - New dude - Create seperate dumps per backups
# 2010/09/03 - Ops dude - Add set -e to make sure backup actually fails
# 2010/09/15 - Dev dude - Compress backups to save space
```

# Looks good to me!



KubeCon



CloudNativeCon

China 2018

```
# important don't remove!!!
set -e

echo "Started backup"
start=`date +%s`
timestamp=$(date+%s)

for table in $(echo "use jobsite_db; show tables;" | mysql -uroot -psupersecretpassword); do
    echo "Backup up table: ${table}"
    mysqldump jobsite_db ${table} | gzip > /mnt/backup/${timestamp}/${table}.sql.gz
    echo "Finished backing up table: ${table}"
done

end=`date +%s`
runtime=$((end-start))
echo "Backup successful!!! Finished in ${runtime} seconds!"
```

# This is fine



KubeCon



CloudNativeCon

China 2018

No table locking

Need to restore the tables in the reverse order

Need to write a restore script

Only alerts when something  
went wrong



# Things that email are good at



KubeCon



CloudNativeCon

China 2018

1

Humans sending a message to another human  
who already knows the original human

2

Saving articles for yourself that you plan on reading later



# Things that email are not so good at



KubeCon



CloudNativeCon

China 2018



- ❌ Monitoring alerts
- ❌ Computers contacting humans
- ❌ Alerting the fire brigade

# Wow, it actually worked



KubeCon



CloudNativeCon

China 2018

Restore worked perfectly first go

Site started up

Site actually worked



# ‘Worked’ is a loaded term



KubeCon



CloudNativeCon

China 2018

The exit code of my restore script was 0

Stuff was on the website

Stuff was in the database

There were actually only 8 jobs online

# OK, let's go back in time



KubeCon



CloudNativeCon

China 2018

git blame: "initial commit from svn"

"Maybe it was just a **bad backup**? Let's look at yesterday's backup."



# Hmm...



KubeCon



CloudNativeCon

China 2018

```
$ du -hs /mnt/backups/*/jobs.tar.gz
```

```
1.0M    1539761034/jobs.tar.gz
```

```
500.3M  1539674634/jobs.tar.gz
```

```
183.5M  1539588234/jobs.tar.gz
```

```
302.3M  1539501834/jobs.tar.gz
```

```
10.7M   1539415434/jobs.tar.gz
```

```
485.9M  1539329034/jobs.tar.gz
```

```
56.3M   1539242634/jobs.tar.gz
```

```
152.3M  1539156234/jobs.tar.gz
```

```
501.3M  1539069834/jobs.tar.gz
```

# Let's take a closer look



KubeCon



CloudNativeCon

China 2018

```
# important don't remove!!!
set -e

echo "Started backup"
start=`date +%s`
timestamp=$(date+%s)

for table in $(echo "use jobsite_db; show tables;" | mysql -uroot -psupersecretpassword); do
    echo "Backup up table: ${table}"
    mysqldump jobsite_db ${table} | gzip > /mnt/backup/${timestamp}/${table}.sql.gz
    echo "Finished backing up table: ${table}"
done

end=`date +%s`
runtime=$((end-start))
echo "Backup successful!!! Finished in ${runtime} seconds!"
```

# Changelog



KubeCon



CloudNativeCon

China 2018

```
#!/bin/bash
# Author: Some dude
#
# Changelog
# 1995/05/23 - Some dude - Added mysql backups for important company database
# 1995/05/27 - Some dude - Added debug logging and log backup duration
# 1998/02/13 - New dude - Create seperate dumps per backups
# 2010/09/03 - Ops dude - Add set -e to make sure backup actually fails
# 2010/09/15 - Dev dude - Compress backups to save space
```

# But, you were supposed to help me?



KubeCon



CloudNativeCon

China 2018

It's OK if every part of the script failed **except for the last line**

**default**

Do you care if **any** of them failed or just the **last one**?

```
set -e
```

OK, but do you care if **parts of each line** failed or just the **last one**?

```
set -o pipefail
```





KubeCon



CloudNativeCon

China 2018

“ ”

**Software is easy to use**  
**but hard to run.**

- Me again



KubeCon



CloudNativeCon

China 2018

“ ”

**Data is easy to backup,  
but hard to restore.**

- Always me

# Times have changed



KubeCon



CloudNativeCon

China 2018

OK, this was 1995...

And it was bash

So things have improved right?

# Hello Docker



KubeCon



CloudNativeCon

China 2018

```
$ docker run hello-world
```

```
Hello from Docker!
```

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.  
(amd64)
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.





KubeCon



CloudNativeCon

China 2018



**The Docker daemon**  
**pulled the "hello-world" image**  
**from the Docker Hub.**

- Docker daemon (in haiku)



# What actually happened



KubeCon



CloudNativeCon

China 2018

The docker daemon checked if the image `hello world`  
with the tag `latest` was already downloaded

It already existed, so it actually did **nothing**  
and didn't connect to Docker hub

# What didn't happen



KubeCon



CloudNativeCon

China 2018

The Docker Daemon didn't check that this was the right image  
from the Docker Hub

It didn't download `docker.io/library/hello-world`

It didn't download the latest `latest`

# Nit picking?



KubeCon



CloudNativeCon

China 2018

`latest` is the default tag for the client -

it doesn't have to be **latest**, or **updated** or even **exist** in the registry

If you already have some version of this tag it won't pull a new one

Your `latest` may not be my `latest` or even exist on Docker Hub anymore

# Down the rabbit hole



KubeCon



CloudNativeCon

China 2018

- ➡ 3 servers, 3 rabbitmq nodes running in Docker with image “rabbitmq”
- ➡ All running 2.8.7 at the time
- ➡ Updates happened to the rabbitmq docker image
- ➡ No updates happened to our setup
- ➡ New server was added to the pool
- ➡ Yay, we just upgraded to version 3.0.0!



# Why haven't you patched prod?



KubeCon



CloudNativeCon

China 2018

```
FROM centos:6.6
```

```
COPY . /app
```

# Ahh we forgot to pull!



KubeCon



CloudNativeCon

China 2018

```
# Need to pull first to get "real" latest
# https://github.com/moby/moby/issues/13331
docker pull centos:6.6
docker build -t business-app .
```

# Nope, still no



KubeCon



CloudNativeCon

China 2018

## Minor tags

Additionally, images with minor version tags that correspond to install media are also offered. These images **DO NOT** receive updates as they are intended to match installation iso contents. If you choose to use these images it is highly recommended that you include `RUN yum -y update && yum clean all` in your Dockerfile, or otherwise address any potential security concerns. To use these images, please specify the minor version tag:

For example: `docker pull centos:5.11` or `docker pull centos:6.6`

# We did it!



KubeCon



CloudNativeCon

China 2018

```
FROM centos:6.6
RUN yum -y update && yum clean all
COPY . /app
```

# 1 month later



KubeCon



CloudNativeCon

China 2018

“Why haven’t you patched prod?”



# Hello Docker build cache!



KubeCon



CloudNativeCon

China 2018

The cache for `RUN` instructions isn't invalidated automatically during the next build. The cache for an instruction like `RUN apt-get dist-upgrade -y` will be reused during the next build. The cache for `RUN` instructions can be invalidated by using the `--no-cache` flag, for example `docker build --no-cache`.



# Ahh we forgot to disable cache!



KubeCon



CloudNativeCon

China 2018

```
# Need to pull first to get "real" latest
# https://github.com/moby/moby/issues/13331
docker pull centos:6.6

# --no-cache is needed to make sure we actually patch
# https://docs.docker.com/engine/reference/builder/#run
docker build --no-cache -t business-app .
```

# Docs check



KubeCon



CloudNativeCon

China 2018

The docs actually show the output from an old `hello-world` image

I like to think that this is because they haven't done a `docker pull hello-world` in a while (this is not true though)

1. Test that your installation works by running the simple Docker image, `hello-world`:

```
docker run hello-world

Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
ca4f61b1923c: Pull complete
Digest: sha256:ca0eeb6fb05351dfc8759c20733c91def84cb8007aa89a5bf606bc8b315b9fc7
Status: Downloaded newer image for hello-world:latest

Hello from Docker!
This message shows that your installation appears to be working correctly.
...
```

## Updating Images

---

The default pull policy is `IfNotPresent` which causes the Kubelet to skip pulling an image if it already exists. If you would like to always force a pull, you can do one of the following:

- set the `imagePullPolicy` of the container to `Always`.
- omit the `imagePullPolicy` and use `:latest` as the tag for the image to use.
- omit the `imagePullPolicy` and the tag for the image to use.
- enable the [AlwaysPullImages](#) admission controller.

Note that you should avoid using `:latest` tag, see [Best Practices for Configuration](#) for more information.

# Hello Kubernetes



KubeCon



CloudNativeCon

China 2018

```
kubectl run nginx --image=nginx
```

```
kubectl expose deployment nginx --port=80 --type=LoadBalancer
```

# Show me what happened!



KubeCon



CloudNativeCon

China 2018

```
$ kubectl get deploy nginx -o yaml
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: nginx
```

```
  namespace: kubecon
```

```
spec:
```

```
  replicas: 1
```

```
  strategy:
```

```
    type: RollingUpdate
```

```
    rollingUpdate:
```

```
      maxSurge: 1
```

```
      maxUnavailable: 1
```

```
  template:
```

```
    spec:
```

```
      containers:
```

```
        - image: nginx
```

```
          imagePullPolicy: Always
```

```
          name: nginx
```

```
          resources: {}
```

```
      restartPolicy: Always
```

# Some important questions

Is it OK if all containers are running on the same host?

Is it OK if all containers go down and fail to start again during upgrades?

Should these containers always run? Or just sometimes?

Should I check if the container is healthy?

You asked for 3 replicas. Do you really want that many or do you just like that number?



# Update Strategy



KubeCon



CloudNativeCon

China 2018

```
spec:
  replicas: 1
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
```

# What would you expect?



KubeCon



CloudNativeCon

China 2018

I deploy a new version

The new version is deployed

Once it is ready and healthy, the old one is removed

# What actually happens



KubeCon



CloudNativeCon

China 2018

Deploy new version

New version is deployed

At the same time old version is removed

New version fails

Application goes down



# What you probably want



KubeCon



CloudNativeCon

China 2018

```
spec:
  replicas: 1
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
```

# What you really really want



KubeCon



CloudNativeCon

China 2018

```
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
```

# My application is invincible!



KubeCon



CloudNativeCon

China 2018

Kubernetes will restart it if it crashes

“When deploying a new version I can’t break anything!”



# Kubernetes upgrade time!



KubeCon



CloudNativeCon

China 2018

Kubernetes cluster is upgraded

Rolling restart of nodes

Application fails to start

Kubernetes doesn't care  
and continues to take down all 3 replicas



# You forgot pod disruption budgets!



KubeCon



CloudNativeCon

China 2018

```
apiVersion: policy/v1beta1
kind: PodDisruptionBudget
metadata:
  name: nginx
spec:
  maxUnavailable: 1
  selector:
    matchLabels:
      app: nginx
```

# Kubernetes node crashes



KubeCon



CloudNativeCon

China 2018

Kubernetes node crashes

All 3 replicas were running on that node

Kubernetes waits 5 minutes

Kubernetes starts the 3 replicas again



# You forgot pod anti-affinity!



KubeCon



CloudNativeCon

China 2018

```
spec:
  affinity:
    podAntiAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        - labelSelector:
            matchExpressions:
              - key: app
                operator: In
                values:
                  - nginx
            topologyKey: kubernetes.io/hostname
```

# Kubernetes is happy now



KubeCon



CloudNativeCon

China 2018

Application is running great

Other teams are jealous

They too want the Kubernetes

They deploy application with this  
`resources` thing

Your application goes down





# What happened

No resource requests or limit were being set

All containers were “best effort”

Every container got an equal share of the available resources

New application was deployed with resource requests and limits

It decided to use them

No resources leftover for your application

# You forgot resources!



KubeCon



CloudNativeCon

China 2018

```
resources:
```

```
  limits:
```

```
    cpu: 1000m
```

```
    memory: 500Mi
```

```
  requests:
```

```
    cpu: 1000m
```

```
    memory: 500Mi
```



# Oh, and maybe quotas and defaults



KubeCon



CloudNativeCon

China 2018

## Resource Quotas

<https://kubernetes.io/docs/tasks/administer-cluster/manage-resources/quota-memory-cpu-namespace/>

## Default resource limits

<https://kubernetes.io/docs/tasks/administer-cluster/manage-resources/cpu-default-namespace/>

<https://kubernetes.io/docs/tasks/administer-cluster/manage-resources/memory-default-namespace/>

# Something is wrong with my app



KubeCon



CloudNativeCon

China 2018

My code has a problem

One container isn't feeling so good

Requests keep getting sent  
even though it is clearly broken



# You forgot probes!



KubeCon



CloudNativeCon

China 2018

readinessProbe:

httpGet:

path: /health

port: 8080

livenessProbe:

httpGet:

path: /health

port: 8080

# Production ready checklist (save this)



KubeCon



CloudNativeCon

China 2018

- Resource Requests/Limits and Quotas
- Liveness and Readiness Probes
- Pod disruption budgets
- Anti-affinity rules
- Upgrade strategy
- Logging, Monitoring and Alerting
- A way to restore your application if Kubernetes gets deleted
  - Version controlled deployments
  - Secret persistence

# Why?



KubeCon



CloudNativeCon

China 2018

First impressions matter

Convenience is king

Not everyone cares about HA and performance

Backwards compatibility

# Can we make things better?



KubeCon



CloudNativeCon

China 2018

- ➡ Simpler documentation with 'production ready' examples?
- ➡ Stricter defaults that prioritise safety over compatibility?
- ➡ 'Production mode' where you are forced to always set these values?
- ➡ Include these features in the default helm charts templates?



KubeCon



CloudNativeCon

China 2018



**Deploying to production is *easy*.**

**Running tests locally is *hard*.**

- Everyone (including me)



# Lope this way



KubeCon



CloudNativeCon

China 2018

<https://github.com/Crazybus/lope>

- Automatically mount current working directory
- Forward environment variables
- Forward your SSH agent into the container
- Forward docker socket and automatically add a docker client



**KubeCon**



**CloudNativeCon**

China 2018

