



KubeCon



CloudNativeCon

China 2018

# Navigating the Distributed Systems Execution Maze with OpenTracing

Ashlie Martinez

Ilya Kislenko (on behalf of Julio López)



# Why Do We Need Distributed Tracing?




KubeCon



CloudNativeCon

China 2018

- K10 (Data Management for Cloud-Native Applications) has 13 microservices
- Don't want to redesign existing K10 microservices just for tracing
- Logging exposes some information, but cannot show time spent in services
- Require tracing library that can be incrementally added



# Sample App: Image Gallery



KubeCon




CloudNativeCon

China 2018



Latency < 100 ms



# Sample App: Image Gallery

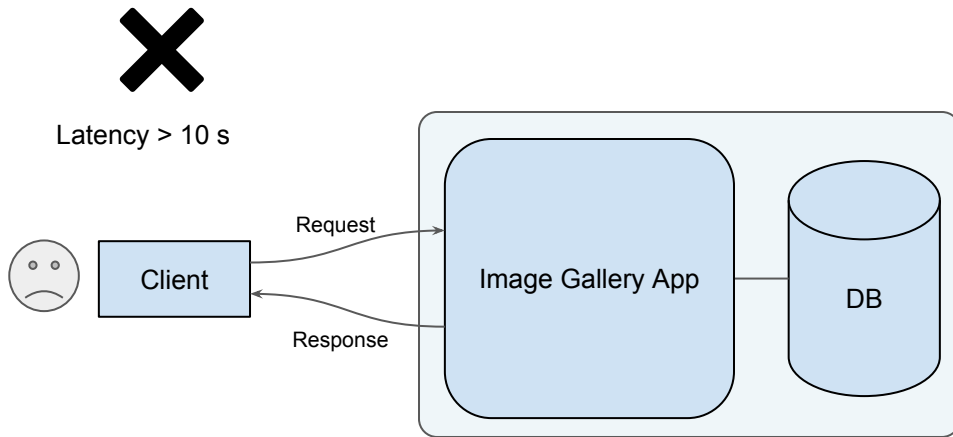


KubeCon



CloudNativeCon

China 2018



# Sample App: Image Gallery

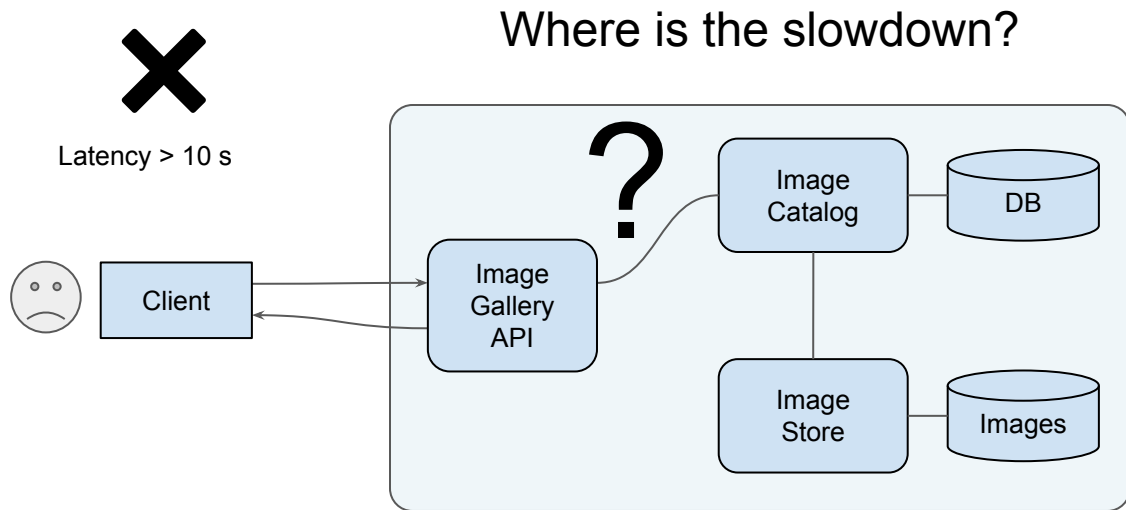


KubeCon



CloudNativeCon

China 2018



It becomes much harder to troubleshoot and debug a set of intertwined distributed microservices

# Distributed Tracing



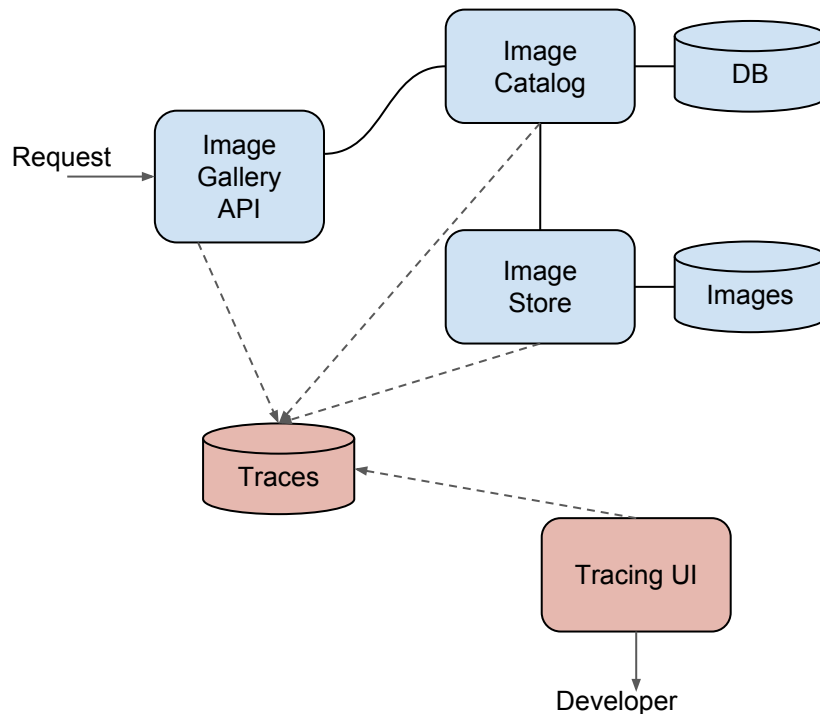
KubeCon



CloudNativeCon

China 2018

- Automatically aggregate traces for a request from multiple microservices
- Highlights the execution path of client requests in the backend
- Can help pinpoint where failures occur and what causes poor performance
- Most libraries support probabilistic sampling of requests
- Complements distributed log and metrics collection tools



# Leveraging Distributed Tracing: Outline



Steps for tracing applications using REST-based Go microservices running in K8s

- How does distributed request tracing work?
- Our choice for tracing: OpenTracing and Jaeger
- Installing Jaeger in a K8s cluster
- Instrumenting Go microservices using OpenTracing SDKs and APIs
  - Request IDs and spans
- Dealing with services external to the application: Cloud Providers and K8s API

# How Does Distributed Tracing Work?

General approach:

- Instrument parts of services with tracing framework to record information
  - Can instrument any part of service, though request level gives reasonable visibility into system
- Configure services to send tracing data to a central database for display
- Database correlates traces from different services for the same request
- Use separate UI to display and search tracing data



# OpenTracing & Jaeger



KubeCon



CloudNativeCon

China 2018

## Open Tracing:

- CNCF distributed tracing library for Go, C#, Java, and other languages
- Instrument existing code with OpenTracing calls to collect tracing information

## Jaeger:

- CNCF UI for visualizing and searching tracing data
- Uses coalesced tracing data stored in a database like Cassandra
- Deployable via helm chart and K8s yaml

Other tracing options: Zipkin, Google OpenCensus



KubeCon



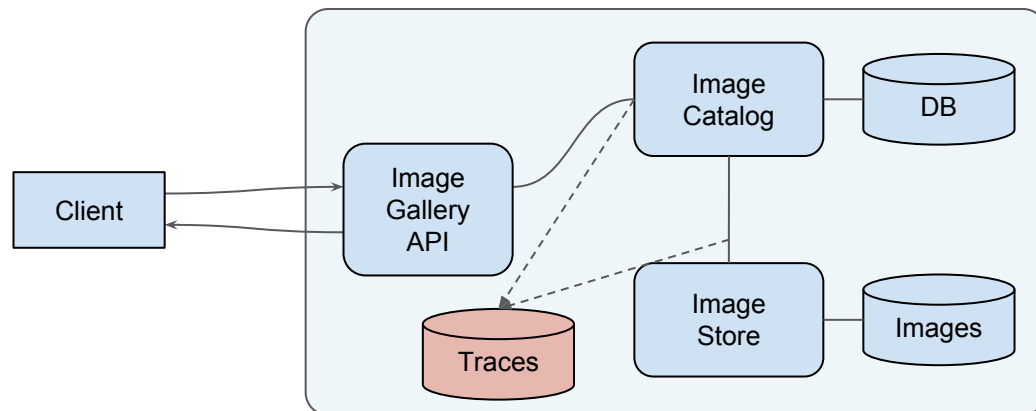
CloudNativeCon

China 2018

# Instrumenting Image Gallery App: Part 1

## Instrumenting the Image Gallery API microservice with custom Middleware

```
func Middleware(next http.Handler) http.Handler {  
    // requests that go through it.  
    return nethttp.Middleware(opentracing.GlobalTracer(),  
        next,  
        nethttp.OperationNameFunc(func(r *http.Request) string {  
            return "HTTP " + r.Method + " " + r.URL.String()  
        })))  
}
```



# OpenTracing Go SDK



KubeCon



CloudNativeCon

China 2018

- Each trace collected by a single service is called a “span”
  - Spans can be nested to show one service calling another
- OpenTracing leverages Go’s Context object to carry info about traces
  - Code being traced must propagate Context to be traced
- Information like HTTP status codes or request IDs can be added to traces
  - Allows developers to get more information about the state of the system for that trace
  - Can help the developer associate a specific trace with other debug information like logs
- Deploy Jaeger with in-memory database via Helm chart



KubeCon



CloudNativeCon

China 2018

# Instrumenting Image Gallery App: Part 1

And we got 1 lonely trace

✓ metadata: HTTP GET /v0/images/412ac325-e0b0-11e8-999a-72fafaa730c4

🔍 Search... ⬆ ⬇ ⬅ View Options ▾

Trace Start: November 4, 2018 8:07 PM | Duration: 177.21ms | Services: 1 | Depth: 1 | Total Spans: 1



Service & Operation

⌵ > ⌵ ⌵

0ms 44.3ms 88.61ms 132.91ms 177.21ms

metadata HTTP GET /v0/images/412ac325-e0b0-11e8-999a-72fafaa730c4

HTTP GET /v0/images/412ac325-e0b0-11e8-999a-72fafaa730c4

Service: metadata | Duration: 177.21ms | Start Time: 0ms

> Tags: sampler.type = const | sampler.param = true | span.kind = server | http.method = GET | http.url = /v0/images/412ac325-e0b0-11e8-999a-72fafaa730c4 | component = net/http | http.status\_code = 200

> Process: client-uuid = 45f7210cc2a1d6ec | hostname = metadata-56d4c88946-dfvd8 | ip = 10.4.0.41 | jaeger.version = Go-2.14.0

SpanID: 6a202491b1144b39



KubeCon



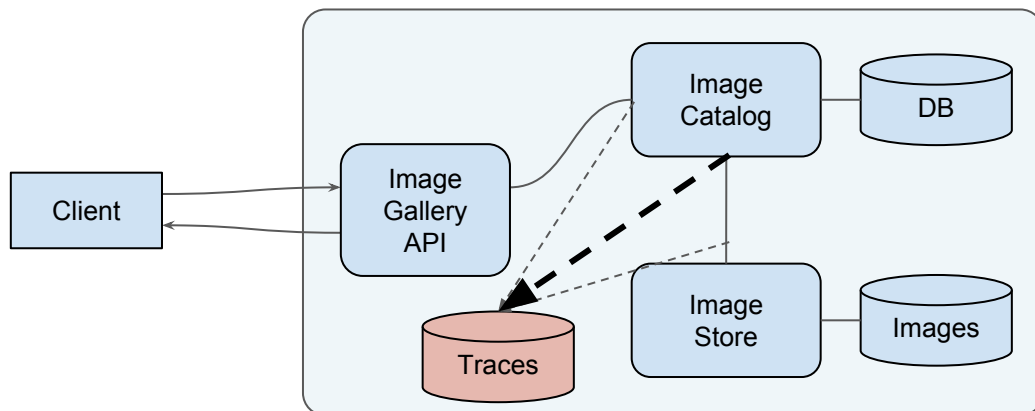
CloudNativeCon

China 2018

# Instrumenting Image Gallery App: Part 2

## Instrumenting the Image Catalog and Image Store microservices' outgoing requests

```
func (t *tracingTransport) RoundTrip(r *http.Request) (*http.Response, error) {  
    ctx := r.Context()  
    span, ctx2 := opentracing.StartSpanFromContext(ctx, "HTTP Request")  
    defer span.Finish()  
  
    r.WithContext(ctx2)  
    carrier := opentracing.HTTPHeadersCarrier(r.Header)  
    span.Tracer().Inject(span.Context(), opentracing.HTTPHeaders, carrier)  
  
    resp, err := t.transport.RoundTrip(r)  
  
    return resp, err  
}
```





KubeCon



CloudNativeCon

China 2018

# Instrumenting Image Gallery App: Part 2

Now we can see that metadata is calling store service

✓ metadata: HTTP GET /v0/images/e077e436-e077-11e8-96d2-1e4a64d5967b




Search...



View Options

Trace Start: November 4, 2018 1:23 PM | Duration: 670.07ms | Services: 1 | Depth: 2 | Total Spans: 2



Service & Operation | 0ms | 167.52ms | 335.03ms | 502.55ms | 670.07ms

metadata HTTP GET /v0/images/e077e436-e077-11e8-96d2-1e4a64d5967b

HTTP GET /v0/images/e077e436-e077-11e8-96d2-1e4a64d5967b

Service: metadata | Duration: 670.07ms | Start Time: 0ms

> Tags: sampler.type = const | sampler.param = true | span.kind = server | http.method = GET | http.url = /v0/images/e077e436-e077-11e8-96d2-1e4a64d5967b | component = net/http | http.status\_code = 200

> Process: client-uuid = 4732451527c6ba44 | hostname = metadata-8477dc4cb6-fm65x | ip = 10.4.0.30 | jaeger.version = Go-2.14.0

SpanID: 2f959770afe82cd5

metadata HTTP Request | 28.91ms

HTTP Request

Service: metadata | Duration: 28.91ms | Start Time: 4.53ms

> Tags: span.kind = server | http.method = GET | http.url = http://store:8000/v0/store | http.status\_code = 200

> Process: client-uuid = 4732451527c6ba44 | hostname = metadata-8477dc4cb6-fm65x | ip = 10.4.0.30 | jaeger.version = Go-2.14.0

SpanID: 3f40410995301c40



KubeCon



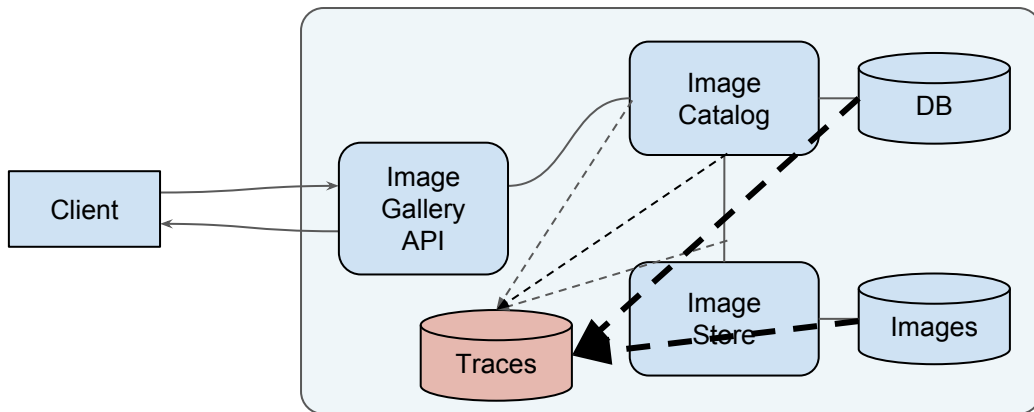
CloudNativeCon

China 2018

# Instrumenting Image Gallery App: Part 3

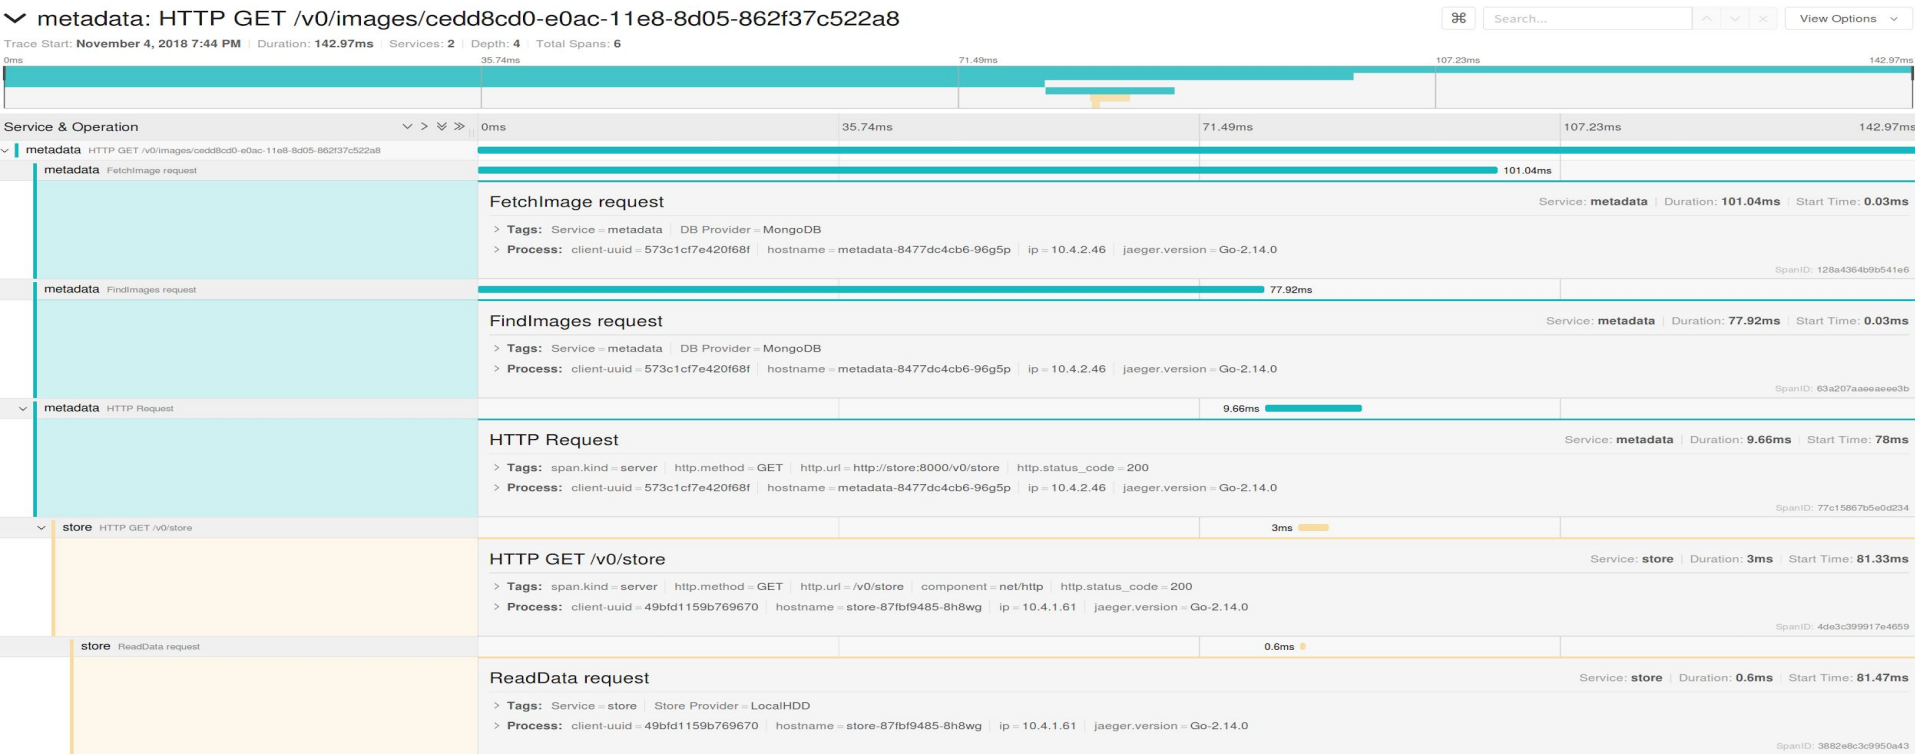
## Internal instrumentation example: Calls into the DB

```
func (s *Mongo) GetAllImages(ctx context.Context) (models.ImageList, error) {  
    span, _ := opentracing.StartSpanFromContext(ctx, "GetAllImages request")  
    defer span.Finish()  
    addSpanTags(span)  
    err := s.Ping()  
    if err != nil {  
        return models.ImageList{}, err  
    }  
    c := s.Conn.DB(dbName).C(collName)  
    imgs := models.ImageList{}  
    return imgs, c.Find(nil).All(&imgs)  
}
```



# Instrumenting Image Gallery App: Part 3

Here we can see everything.





# Discussion



KubeCon



CloudNativeCon

China 2018

General approach: start with coarse tracing and add finer granularity as needed

Trade-offs:

- Pro: Fine granularity and detailed request information
- Con: Additional resource requirements.
  - Request processing in each of the services, and additional network traffic
  - Additional processing and storage requirements for the traces

Other thoughts:

- Service meshes, such as Istio, give you coarse grained tracing
- Complements and overlaps with metrics and logging

# Bringing it all Together



KubeCon



CloudNativeCon

China 2018

Today we discussed

- Using OpenTracing Go SDK to add instrumentation microservices
- Instrumenting calls to other services: DB, cloud provider, K8s API
- Installing Jaeger tracing collector and UI in k8s cluster
- Using Jaeger UI to visualize, analyze and dig into traces

# Final Thoughts

Tracing can give insights into system bottlenecks, but need to balance with time spent adding instrumentation

## Trade-offs:

- Pro: Fine granularity and detailed request information
- Con: Additional resource requirements.
  - Request processing in each of the services, and additional network traffic
  - Additional processing and storage requirements for the traces

## Other thoughts:

- Service meshes, such as Istio, give you coarse grained tracing
- Complements and overlaps with metrics and logging



**KubeCon**



**CloudNativeCon**

China 2018

# Questions?



# Sample App: Image Gallery

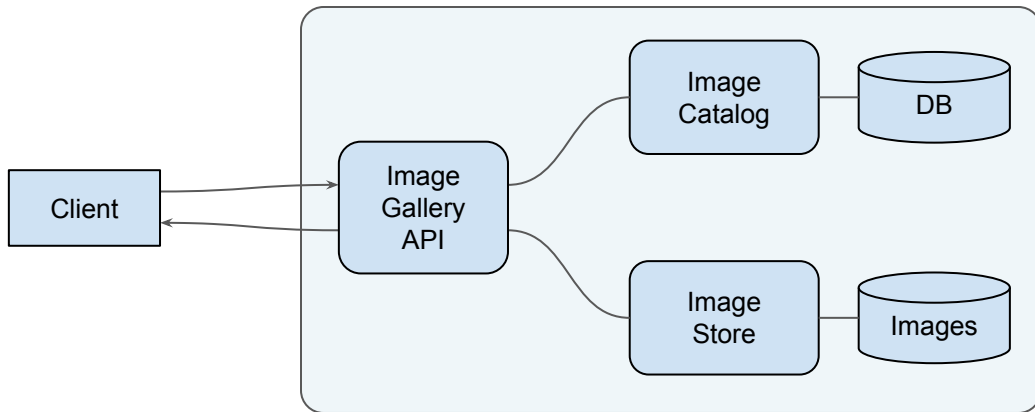


KubeCon



CloudNativeCon

China 2018



# Sample App: Image Gallery

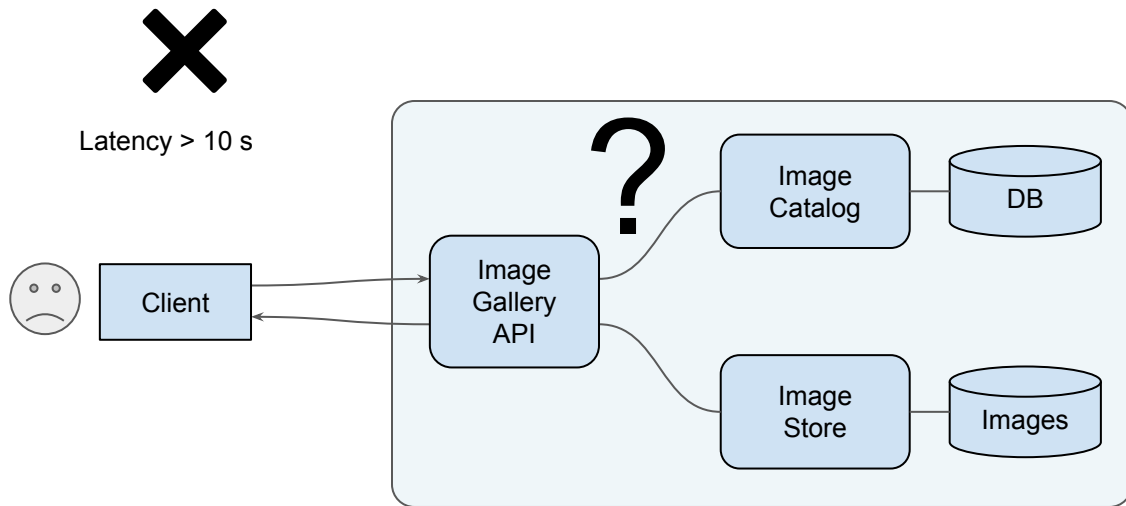


KubeCon



CloudNativeCon

China 2018



Where is the slowdown?