



containerd deep dive

Kubecon Shanghai 2018
Derek McGowan ([Docker](#))
Mike Brown ([IBM](#))










Containerd 1.2 Architecture

clients




containerd



OS



Client-Server



Client

- High level operations using client
- New functionality, interfaces may change (rarely)

Server


- Low level interfaces to resources over GRPC
- Stable API, guaranteed 1.x compatibility



Backend Services

Service

- Provides access to all components
- Low level components wrapped by metadata store
 - Provides labeling
 - Provides namespacing





Namespacing

Metadata

- Support for multiple clients
- Namespaced images
- Namespaced container configurations
- Namespaced snapshots and content
- Namespaced labels

Metadata
(namespaced)

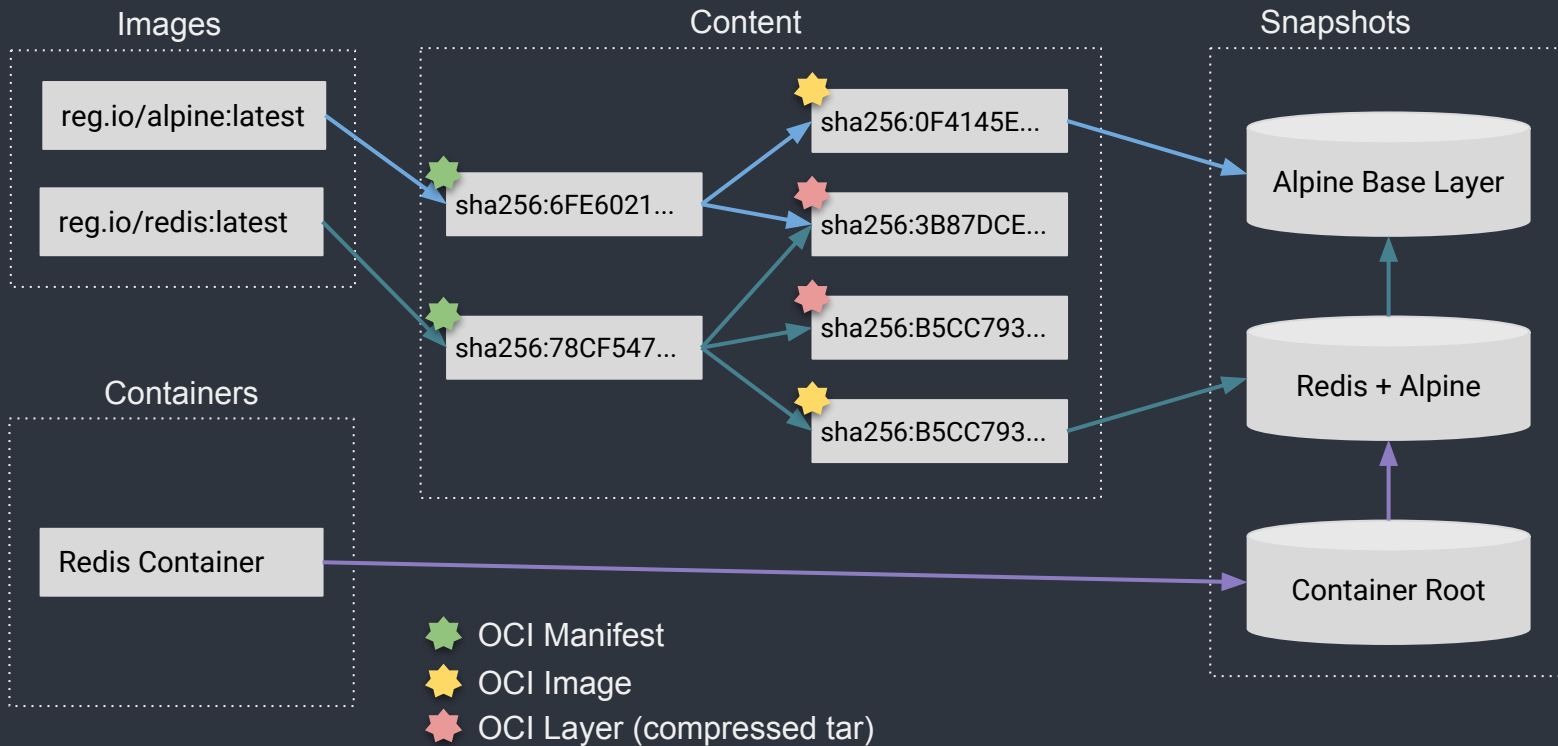
Namespaces

Images

Containers

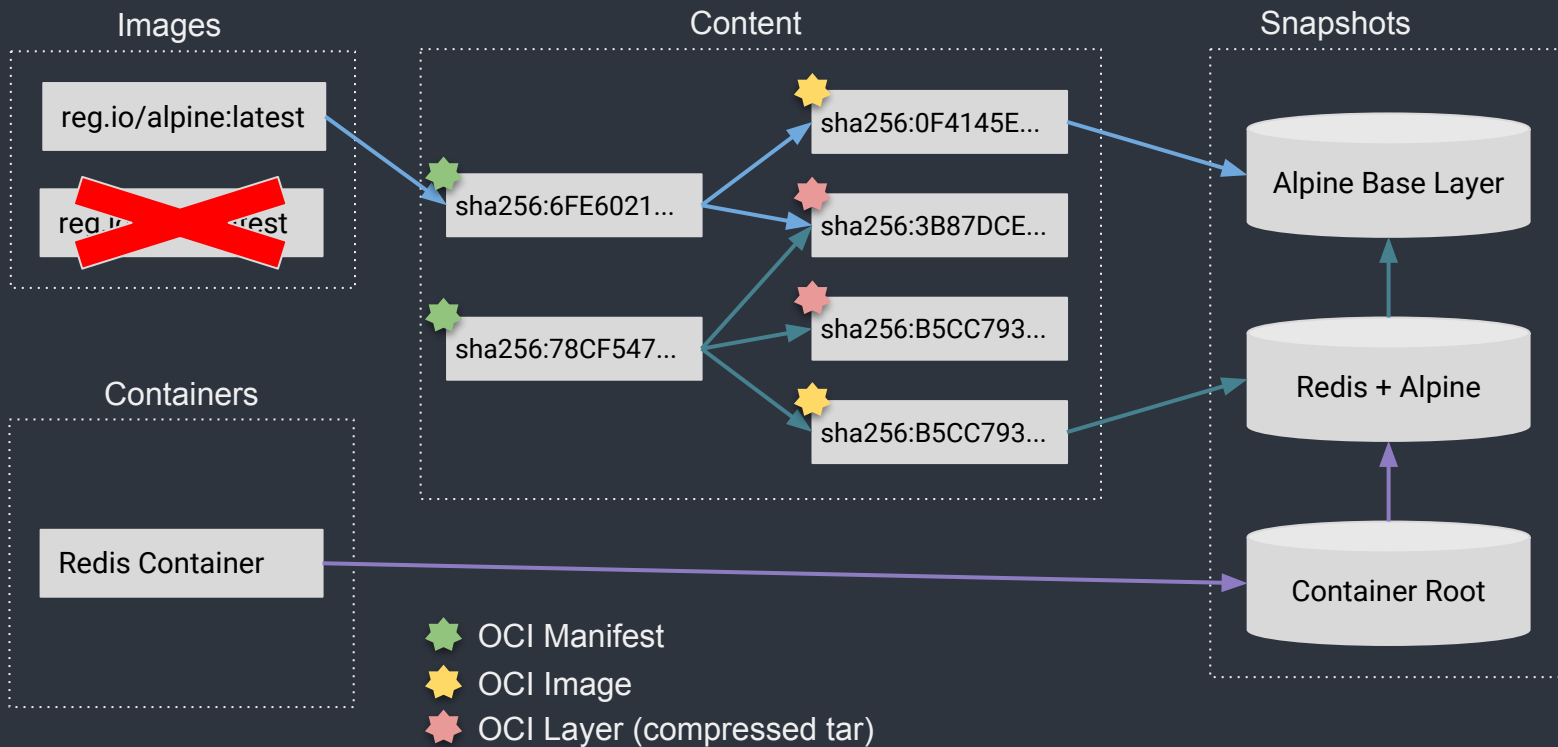


Metadata Deeper Look



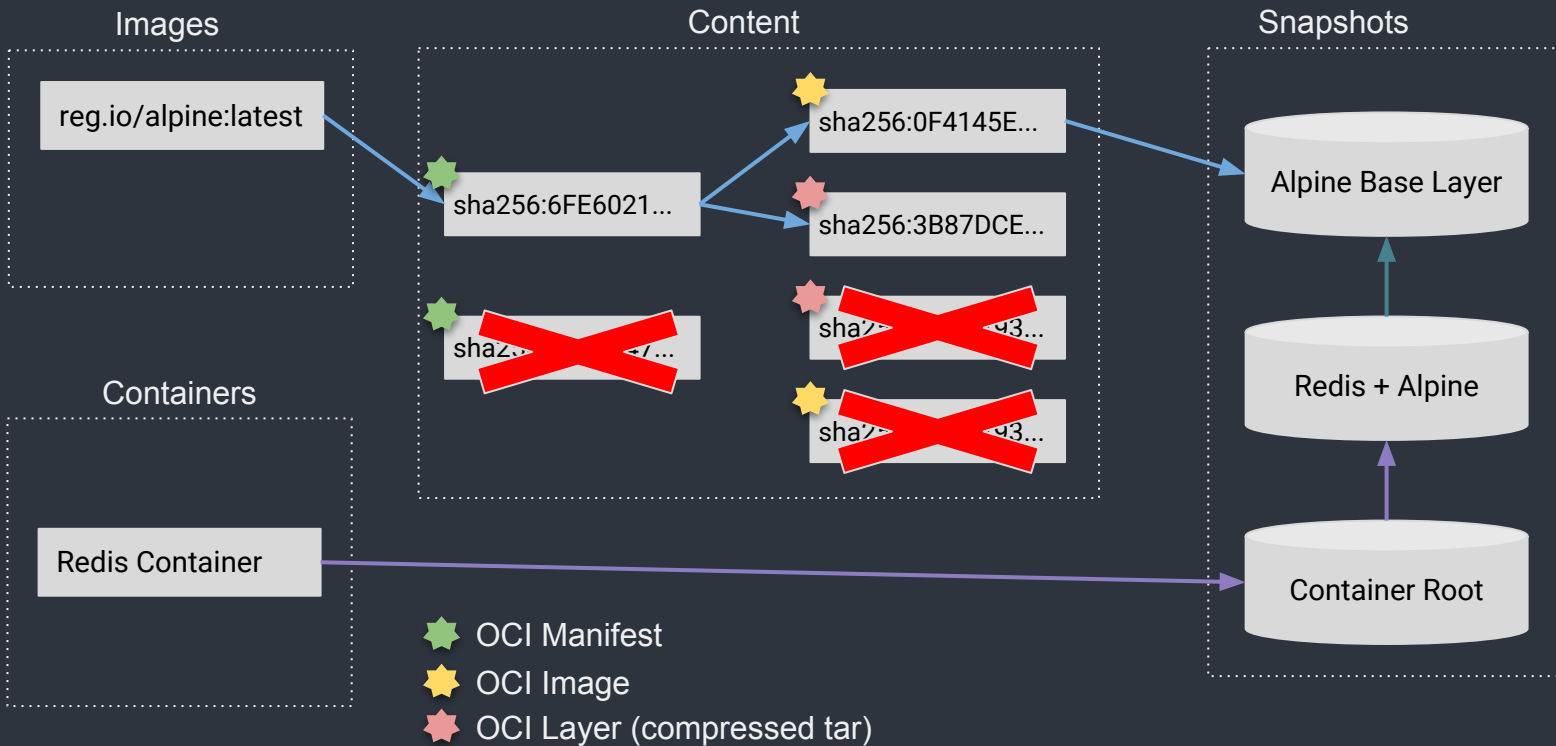


Garbage Collection (delete image)



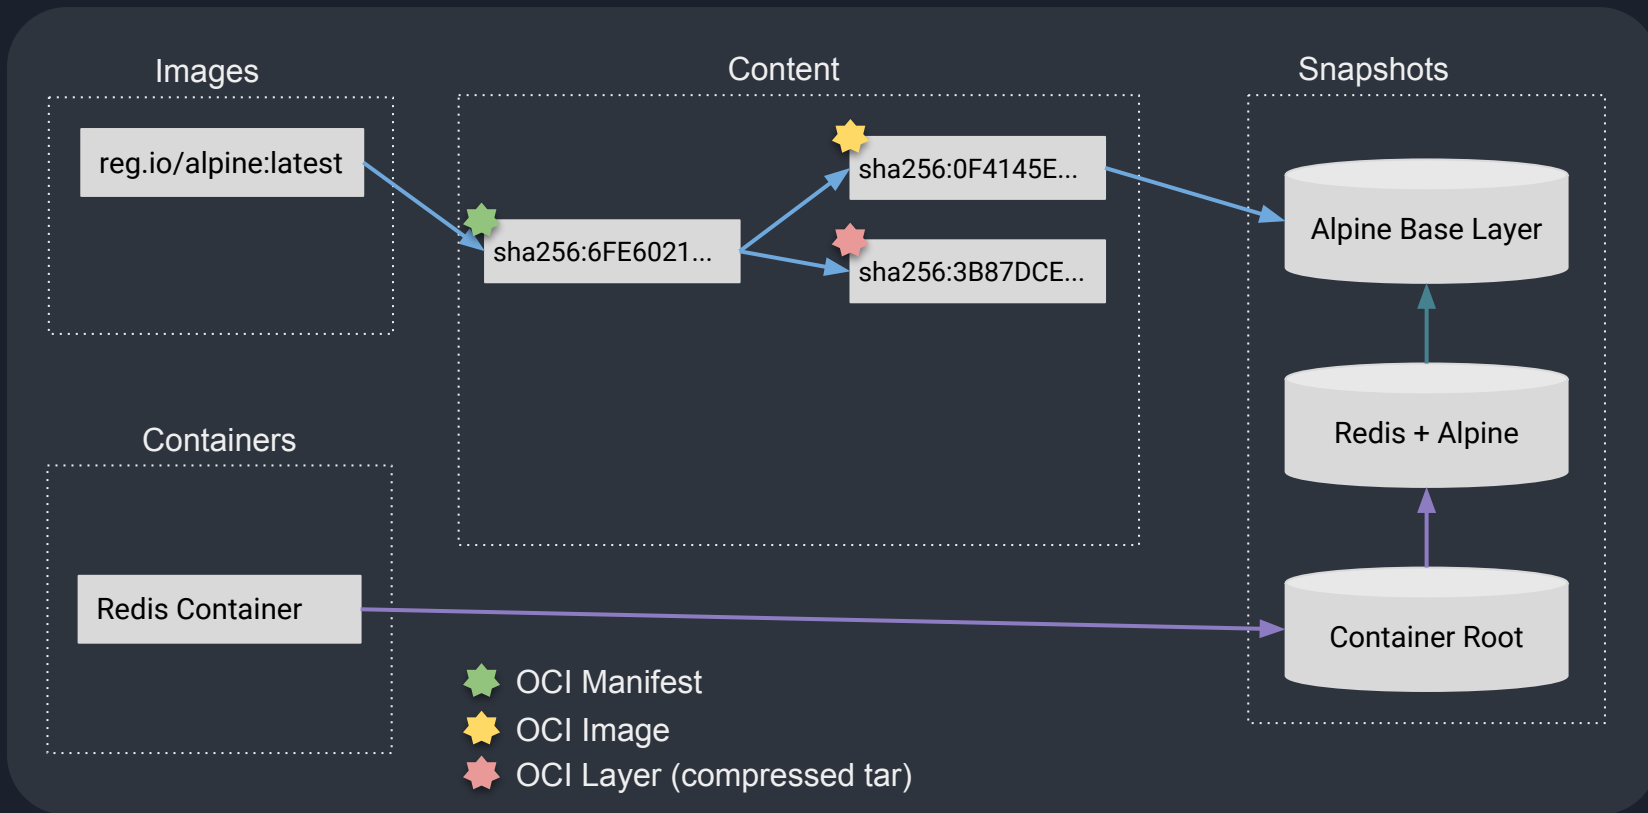


Garbage Collection (running)



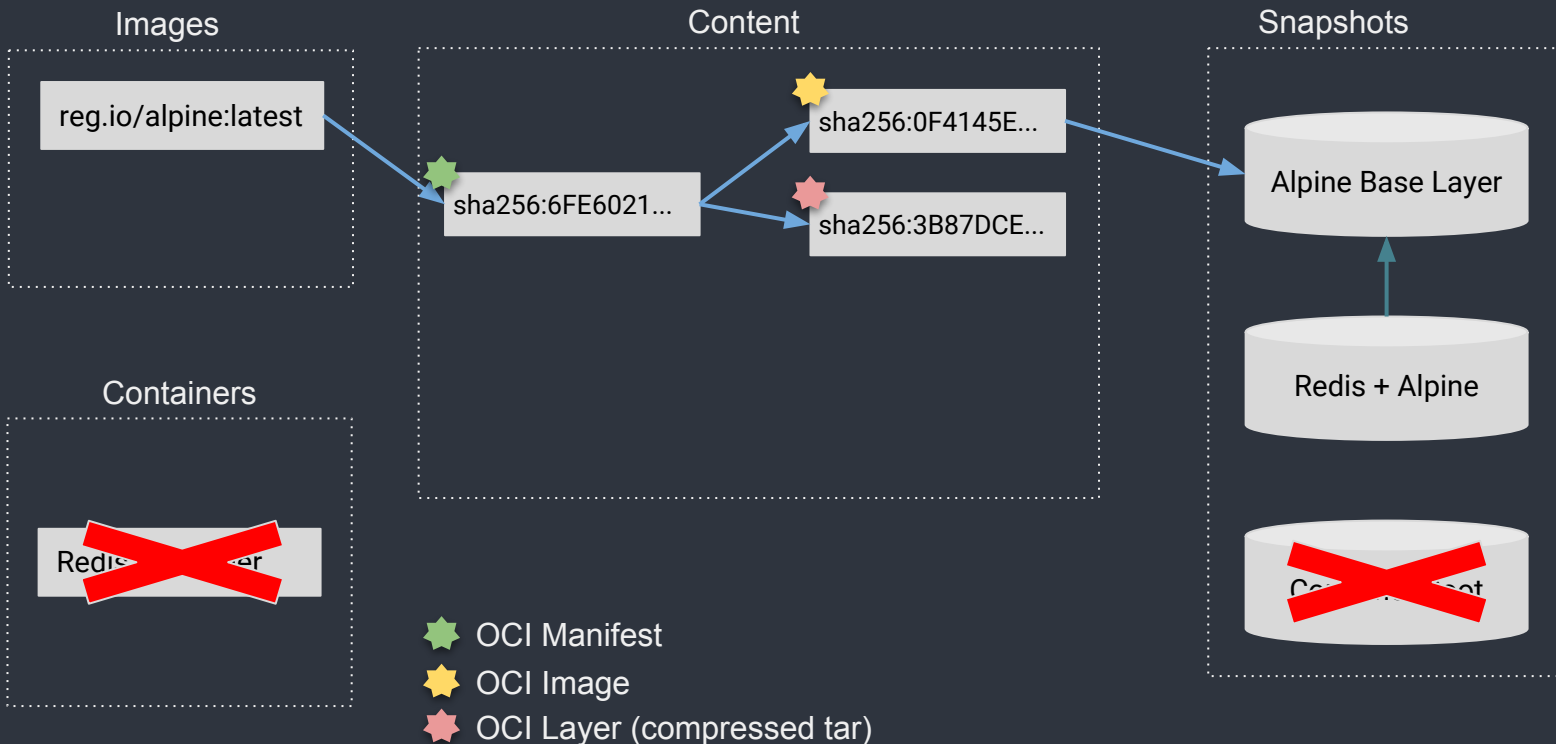


Garbage Collection (after)



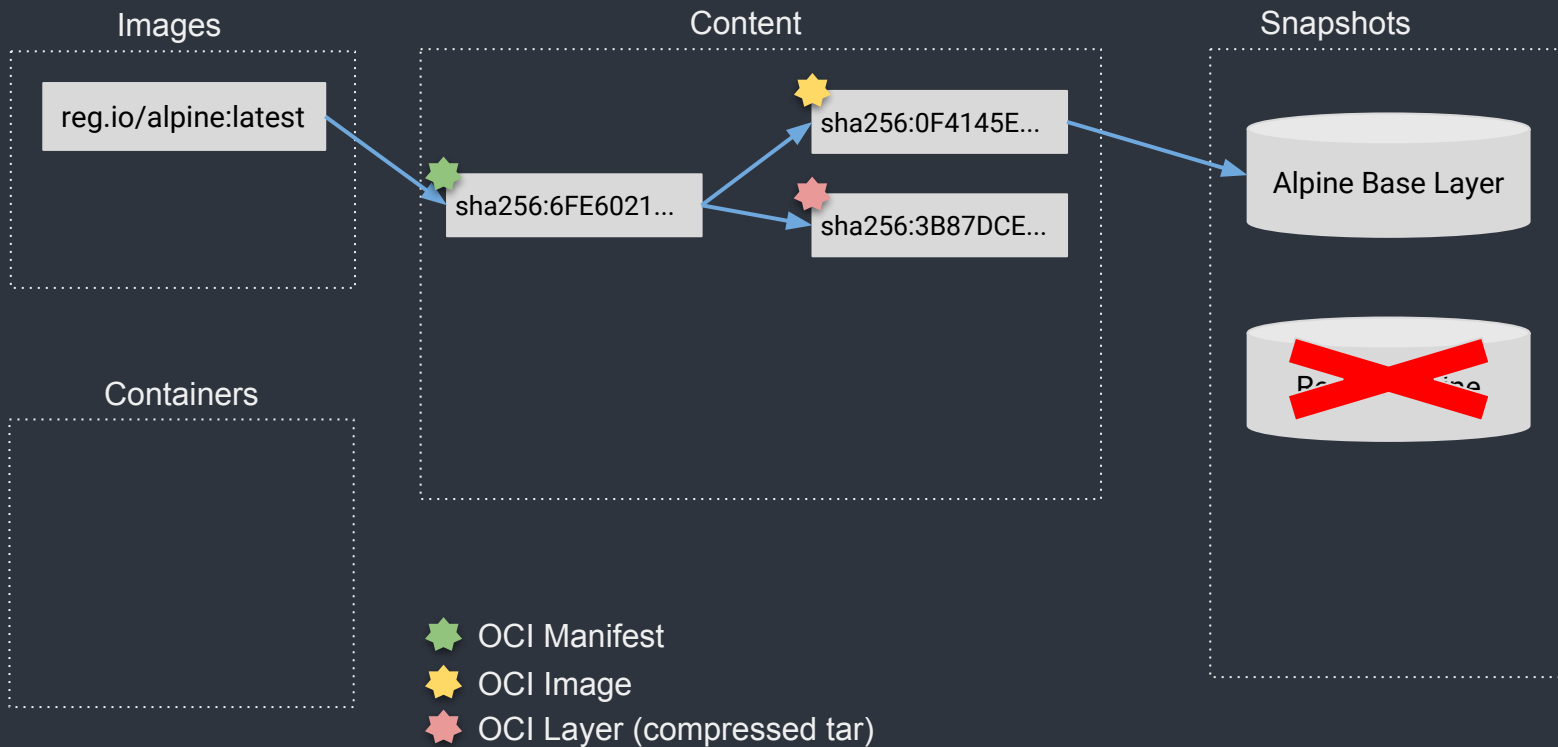


Garbage Collection (delete container)



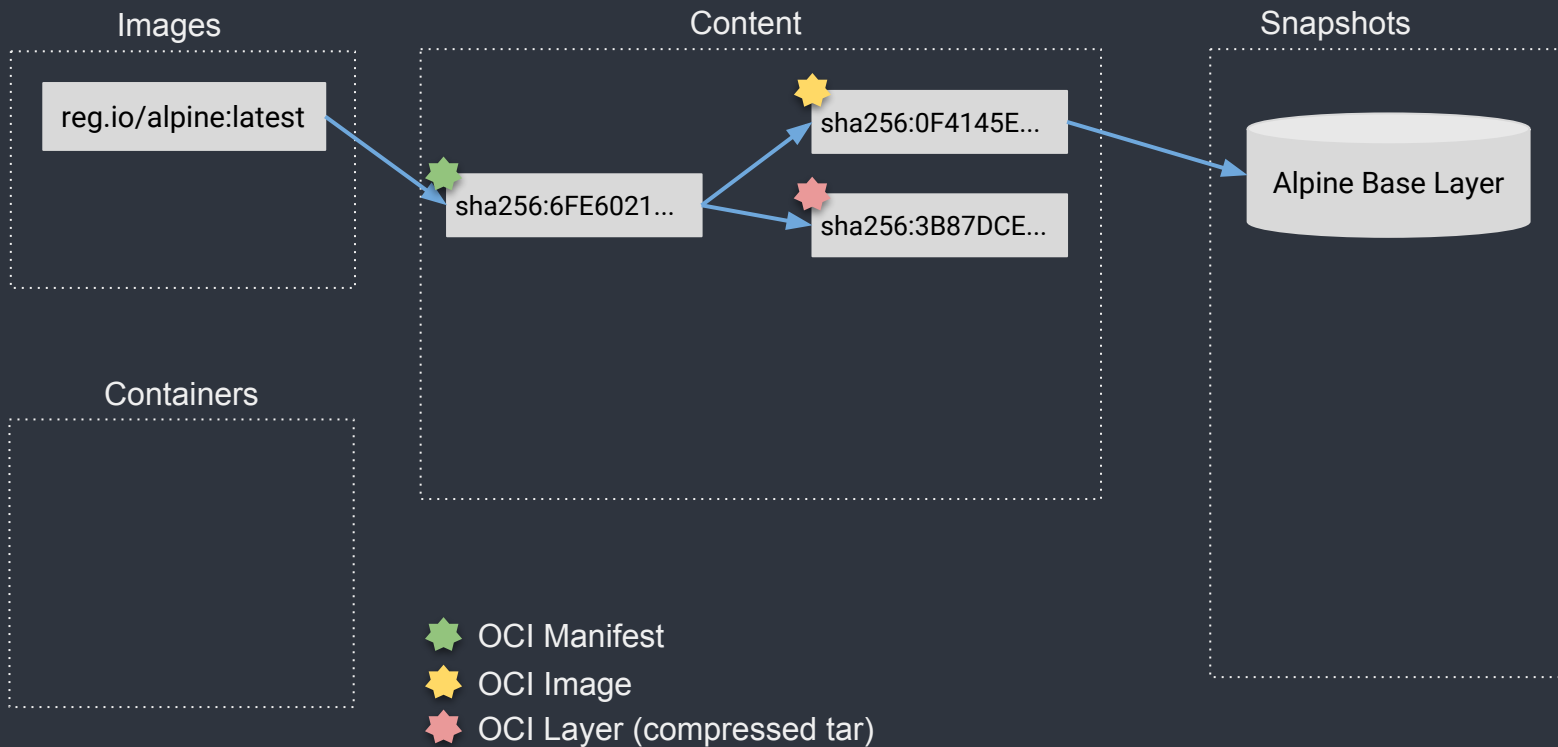


Garbage Collection (running)





Garbage Collection (final)





Plugins (snapshots)

Snapshotter Design

- No data operations
- No mounting
- Immutable snapshots
- Label support
- Enumeration

```
type Snapshotter interface {  
    Stat(Context, string) (Info, error)  
  
    Update(Context, Info, ...string) (Info, error)  
  
    Usage(Context, string) (Usage, error)  
  
    Mounts(Context, string) ([]mount.Mount, error)  
  
    Prepare(Context, string, string, ...Opt) ([]mount.Mount, error)  
  
    View(Context, string, string, ...Opt) ([]mount.Mount, error)  
  
    Commit(Context, string, string, ...Opt) error  
  
    Remove(Context, string) error  
  
    Walk(Context, func(context.Context, Info) error) error  
  
    Close() error  
}
```



Plugins (snapshots)

Proxy Snapshotter

1. Build as an external plugin
2. Configure Containerd to use proxy plugin
(/etc/containerd/config.toml)

2

```
...  
[proxy_plugins]  
  [proxy_plugins.mysnapshotter]  
    type = "snapshot"  
    address = "/tmp/sn.sock"  
...
```

1

```
package main  
  
import(  
    "net"  
    "log"  
    "github.com/containerd/containerd/api/services/snapshots/v1"  
    "github.com/containerd/containerd/contrib/snapshotter"  
)  
  
func main() {  
    rpc := grpc.NewServer()  
    sn := CustomSnapshotter()  
    service := snapshotter.FromSnapshotter(sn)  
    snapshots.RegisterSnapshotsServer(rpc, service)  
  
    // Listen and serve  
    l, err := net.Listen("unix", "/tmp/sn.sock")  
    if err != nil {  
        log.Fatalf("error: %v\n", err)  
    }  
    if err := rpc.Serve(l); err != nil {  
        log.Fatalf("error: %v\n", err)  
    }  
}
```



Plugins (runtime)

Shim GRPC API

- Low level
- Stats now included
- Stabilized in 1.2

Shim Processes

- At most 1 per container
- Can be shared (id passed to every request)

```
service Task {  
  rpc State(StateRequest) returns (StateResponse);  
  rpc Create(CreateTaskRequest) returns (CreateTaskResponse);  
  rpc Start(StartRequest) returns (StartResponse);  
  rpc Delete(DeleteRequest) returns (DeleteResponse);  
  rpc Pids(PidsRequest) returns (PidsResponse);  
  rpc Pause(PauseRequest) returns (google.protobuf.Empty);  
  rpc Resume(ResumeRequest) returns (google.protobuf.Empty);  
  rpc Checkpoint(CheckpointTaskRequest) returns  
(google.protobuf.Empty);  
  rpc Kill(KillRequest) returns (google.protobuf.Empty);  
  rpc Exec(ExecProcessRequest) returns (google.protobuf.Empty);  
  rpc ResizePty(ResizePtyRequest) returns (google.protobuf.Empty);  
  rpc CloseIO(CloseIORequest) returns (google.protobuf.Empty);  
  rpc Update(UpdateTaskRequest) returns (google.protobuf.Empty);  
  rpc Wait(WaitRequest) returns (WaitResponse);  
  rpc Stats(StatsRequest) returns (StatsResponse);  
  rpc Connect(ConnectRequest) returns (ConnectResponse);  
  rpc Shutdown(ShutdownRequest) returns (google.protobuf.Empty);  
}
```



Client Extensibility

Use your own services with service options

```
type ServicesOpt
```


- `func WithContainerService(containerService containersapi.ContainersClient) ServicesOpt`
- `func WithContentStore(contentStore content.Store) ServicesOpt`
- `func WithDiffService(diffService diff.DiffClient) ServicesOpt`
- `func WithEventService(eventService EventService) ServicesOpt`
- `func WithImageService(imageService imagesapi.ImagesClient) ServicesOpt`
- `func WithLeasesService(leasesService leases.Manager) ServicesOpt`
- `func WithNamespaceService(namespaceService namespacesapi.NamespacesClient) ServicesOpt`
- `func WithSnapshotters(snapshotters map[string]snapshots.Snapshotter) ServicesOpt`
- `func WithTaskService(taskService tasks.TasksClient) ServicesOpt`

Customize push and pull with remote options

- `func WithResolver(resolver remotes.Resolver) RemoteOpt`




Flows (pull)



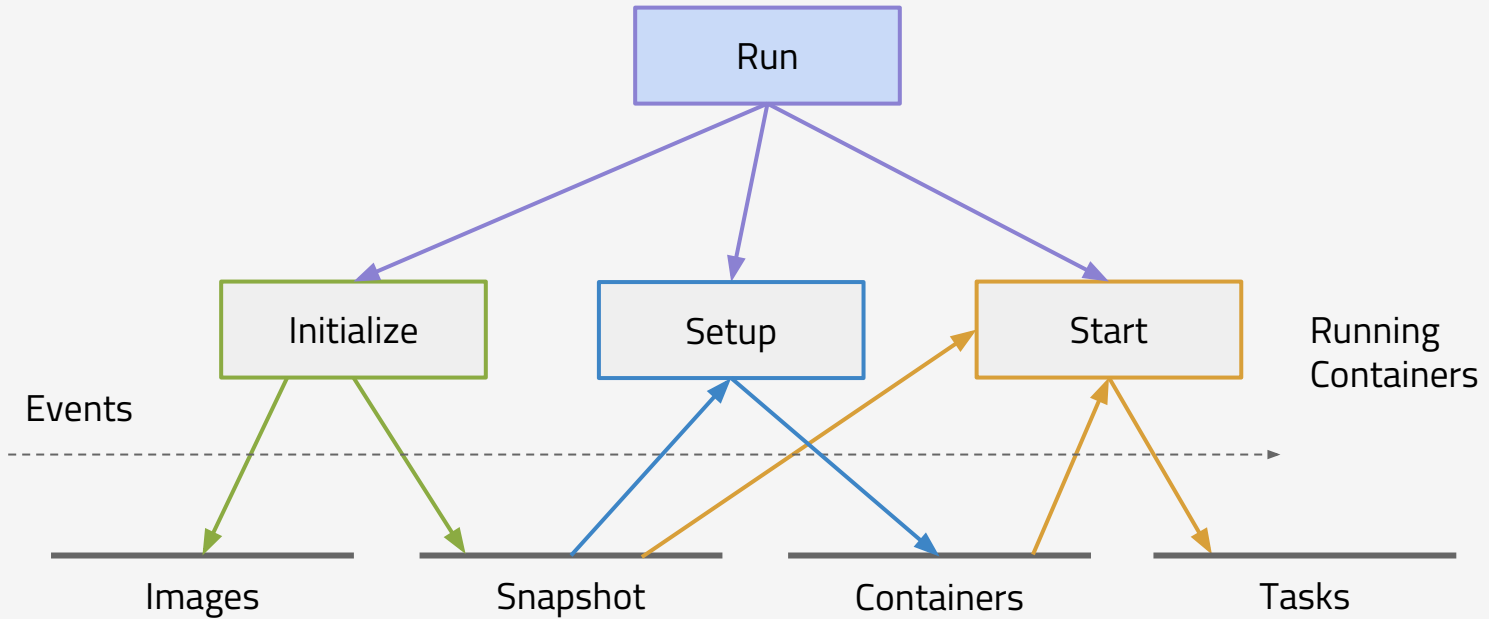


Flow (push)



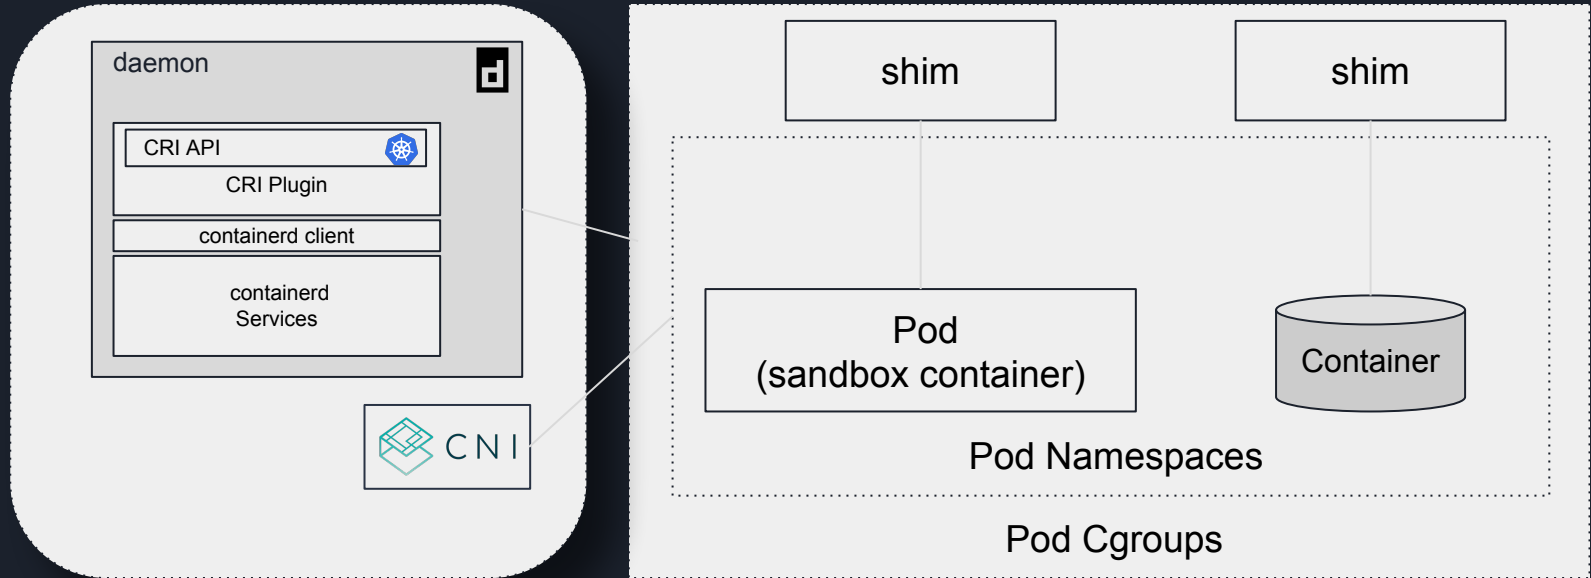


Flow (container run)

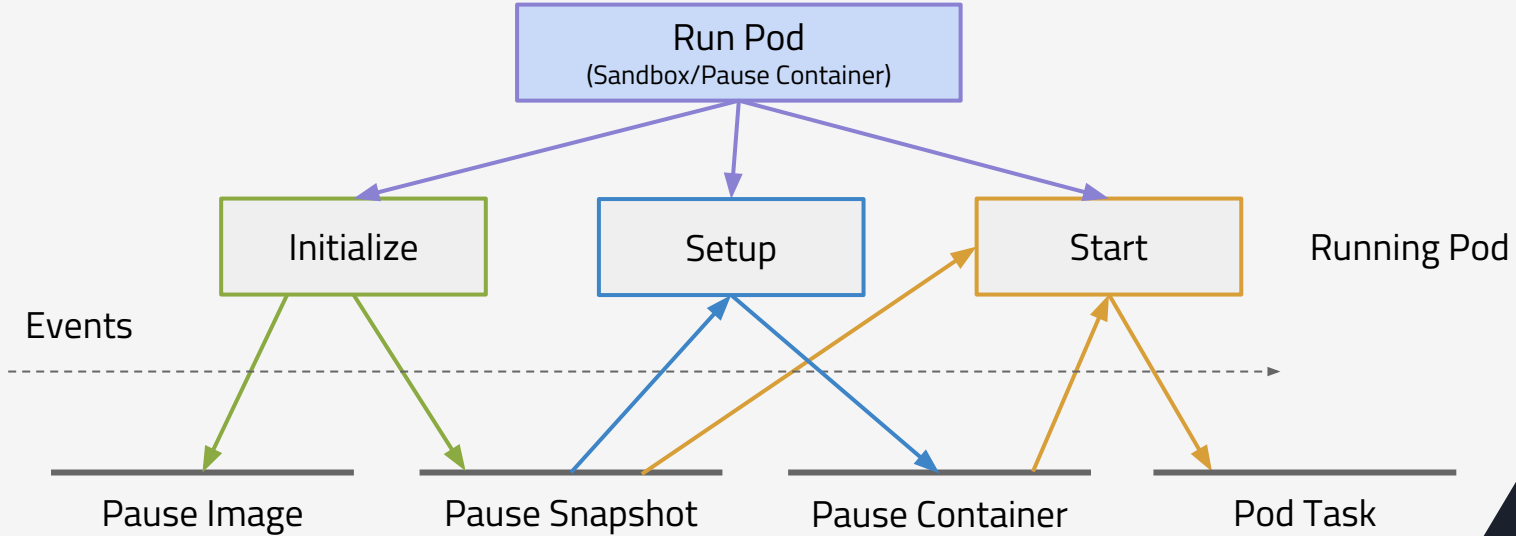





CRI - Pod Namespaces



CRI pod Flow (run pod)



CRI pod Flow (run container)



Demo

