

kubernetes configuration customization

a composable, template-free approach

Nov 2018 kubecon

上海

jeff regan
jregan@google
monopole@github

Kubernetes 配置的用户定制

声明式的途径

Nov 2018 kubecon

上海

jeff regan
jregan@google
monopole@github

~60

k8s configuration tools
k8s的配置工具

recent -
automation broker

#	Project	URL
2		
3	Helm	https://github.com/kubernetes/helm
4	oc	https://docs.openshift.com/online/dev_guide/application_lifecycle/new_app.htm
5	Kompose	https://github.com/kubernetes-incubator/kompose
6	Spread	https://github.com/redislabs/spread
7	ksonnet	https://github.com/ksonnet/ksonnet
8	kubecfg	https://github.com/ksonnet/kubecfg
9	Draft	https://github.com/Azure/draft
10	jsonnet	https://github.com/google/jsonnet
11	Kapitan	https://github.com/deepmind/kapitan
12	Konfd	https://github.com/kelseyhightower/konfd
13	ktmpl	https://github.com/jimmycuadra/ktmpl
14	fabric8 client	https://github.com/fabric8io/kubernetes-client
15	kubegen	https://github.com/errordeveloper/kubegen
16	kenv	https://github.com/thisendout/kenv
17	Ansible	https://github.com/ansible/ansible-kubernetes-modules
18	Puppet	https://forge.puppet.com/qarethr/kubernetes/readme
19	KPM	https://github.com/coreos/kpm
20	Nulecule	https://github.com/projectatomic/nulecule
21	Kedge	https://github.com/kedgeproject/kedge
22	OpenCompose	https://github.com/redhat-developer/opencompose
23	Chartify	https://github.com/appscode/chartify
24	Podex	https://github.com/kubernetes/contrib/tree/master/podex
25	k8sec	https://github.com/dtan4/k8sec
26	kb80r	https://github.com/UKHomeOffice/kb80r
27	k8s-kotlin-dsl	https://github.com/fkorkov/k8s-kotlin-dsl
28	KY	https://github.com/stellarservice/ky
29	kploy	https://github.com/kubernauts/kploy
30	kdeploy	https://github.com/flexiant/kdeploy
31	kubernetes-deploy	https://github.com/Shopify/kubernetes-deploy
32	generator-kubegen	https://github.com/sesispla/generator-kubegen
33	k8comp	https://github.com/cststack/k8comp
34	kontemplate	https://github.com/tazjin/kontemplate
35	kexpand	https://github.com/kopeio/kexpand
36	Forge	https://github.com/datawire/forge/
37	Psykube	https://github.com/CommercialTribe/psykuba
38	Deploymenttizer	https://github.com/InVisionApp/kit-deploymenttizer
39	Broadway	https://github.com/namely/broadway
40	Srvexpand	https://github.com/hortonworks/kubernetes-yarn/tree/master/contrib/srvexpand
41	rok8s-scripts	https://github.com/reactiveops/rok8s-scripts
42	ERB-Hiera	https://github.com/roobert/erb-hiera
43	k82-icl	https://github.com/archipaorg/k8s-icl
44	Compose	https://www.docker.com/kubernetes
45	Deis workflow	https://github.com/deis/workflow
46	OpenShift templates	https://docs.openshift.org/latest/dev_guide/templates.html
47	kube-applier	https://github.com/box/kube-applier

spreadsheet



maintained by Brian Grant

app descriptor



cluster
dashboard



lifecycle
management



Description, maintainer, version, ...





Browse, search, download





Bundling, plus dependencies





What apps are running?
Are they healthy?





Rollouts, rollbacks, upgrades.





Given config - adapt it to *my* needs.

kustomize

Command line tool for k8s customization.
k8s 用户化的命令行工具

Closes several old kubectl issues.
解决了一些kubectl的老问题

Composes with other tools.
可以和别的工具一起使用



```
$ kustomize build helloWorld | \
    kubectl apply -f -
```

```
$ tree helloWorld
```

```
helloWorld
```

```
├── configMap.yaml  
├── deployment.yaml  
├── kustomization.yaml  
└── service.yaml
```

← normal
← normal
← dropped in
← normal

service.yaml

```
kind: Service
metadata:
  name: wordpress
spec:
  ports:
    - port: 389
  selector:
    app: wordpress
```

kustomization.yaml

```
resources:
- service.yaml

namePrefix: demo-
```

kustomize
build

/dev/stdout

```
kind: Service
metadata:
  name: demo-wordpress
spec:
  ports:
    - port: 389
  selector:
    app: wordpress
```

This is k8s-aware patching.

kustomization.yaml =

operands
(things to include)

operations
(ways to patch operands)

operands

operations

result


service.yaml

deployment.yaml


configMap.yaml

namePrefix:
demo-

YAML Stream



operands *operations* *result*




A *result* can be used
as an *operand*.

kustomize input is **plain kubernetes yaml**.

You can ***kubectl apply*** that yaml without kustomize.

To start customizing, just add a **kustomization.yaml** file.

Use Case #1 Variants 不同环境 (development, staging and production)



Use Case #1 Variants 不同环境 (development, staging and production)


kustomization.yaml

```
resources:  
- r1.yaml  
  r2.yaml  
  ...
```

kustomization.yaml

```
namePrefix: dev-  
bases:  
- ../../base
```


dev



kustomization.yaml

```
namePrefix: staging-  
newTag: qa  
bases:  
- ../../base
```


staging



kustomization.yaml

```
namePrefix: prod-  
newTag: v2.3.8  
bases:  
- ../../base
```

production



Use Case #1 Variants 不同环境

(development, staging and production)

File layout:

kustomization.yaml

```
commonLabels:
  app: wordpress
resources:
- deployment.yaml
- service.yaml
configMapGenerator:
- name: wordpress-map
  files:
  - env.txt
```

service.yaml

```
kind: Service
metadata:
  name: wordpress
spec:
  ports:
  - port: 389
```

deployment.yaml

```
kind: Deployment
metadata:
  name: wordpress
spec:
  replicas: 1
  template: ...
```

```
$ tree wordpress
```

```
wordpress
```

```
├── base
│   ├── kustomization.yaml
│   ├── deployment.yaml
│   ├── env.txt
│   └── service.yaml
└── overlays
    ├── production
    │   ├── kustomization.yaml
    │   ├── replica_count.yaml
    │   └── cpu_count.yaml
    ├── staging
    │   ├── kustomization.yaml
    │   └── ...
    └── dev
        ├── kustomization.yaml
        └── ...
```

Use Case #1 Variants 不同环境 (development, staging and production)

kustomization.yaml

```
namePrefix: prod-
commonLabels:
  variant: prod
commonAnnotations:
  note: I'm Prod!
bases:
- ../../base
patchesStrategicMerge:
- replica_count.yaml
- cpu_count.yaml
```

replica_count.yaml

```
kind: Deployment
metadata:
  name: wordpress
spec:
  replicas: 80
```

cpu_count.yaml

```
kind: Deployment
metadata:
  name: wordpress
spec:
  template:
    spec:
      containers:
      - name: my-container
        resources:
          limits:
            cpu: 7000m
```

```
$ tree wordpress
```

```
wordpress
```

```
├── base
│   ├── kustomization.yaml
│   ├── deployment.yaml
│   ├── env.txt
│   └── service.yaml
└── overlays
    ├── production
    │   ├── kustomization.yaml
    │   ├── replica_count.yaml
    │   └── cpu_count.yaml
    ├── staging
    │   ├── kustomization.yaml
    │   └── ...
    └── dev
        ├── kustomization.yaml
        └── ...
```

Deploy production:

```
$ kustomize build \
  wordpress/overlays/production | \
  kubectl apply -f -
```

```
$ tree wordpress
```

```
wordpress
```

```
├── base
│   ├── kustomization.yaml
│   ├── deployment.yaml
│   ├── env.txt
│   └── service.yaml
└── overlays
    ├── production
    │   ├── kustomization.yaml
    │   ├── replica_count.yaml
    │   └── cpu_count.yaml
    ├── staging
    │   ├── kustomization.yaml
    │   └── ...
    └── dev
        ├── kustomization.yaml
        └── ...
```


Deploy staging:

```
$ kustomize build \
  wordpress/overlays/staging |\
  kubectl apply -f -
```

```
$ tree wordpress
```

```
wordpress
```

```
├── base
│   ├── kustomization.yaml
│   ├── deployment.yaml
│   ├── env.txt
│   └── service.yaml
└── overlays
    ├── production
    │   ├── kustomization.yaml
    │   ├── replica_count.yaml
    │   └── cpu_count.yaml
    ├── staging
    │   ├── kustomization.yaml
    │   └── ...
    └── dev
        ├── kustomization.yaml
        └── ...
```

\$ kustomize build target

- 1 load universal k8s object descriptions
- 2 read kustomization.yaml from target
- 3 kustomize bases (*recurse 2-5*)
- 4 load and/or generate resources
- 5 apply target's kustomization operations
- 6 fix name references
- 7 emit yaml



Use Case #2 Feeding customized names to containers

把用户化的名字放入容器

patch.yaml

kustomization.yaml

```
vars:
- name: MYSQL_SERVICE
  objref:
    kind: Service
    name: mysql
    apiVersion: v1
  fieldref:
    fieldpath: metadata.name
patchesStrategicMerge:
- patch.yaml
```

```
kind: Deployment
metadata:
  name: wordpress
spec:
  template:
    spec:
      initContainers:
      - name: init-command
        image: debian
        command:
        - "curl $(MYSQL_SERVICE)"
      containers:
      - name: wordpress
        env:
        - name: WORDPRESS_DB_HOST
          value: $(MYSQL_SERVICE)
```

 /dev/stdout

```
apiVersion: v1
kind: Deployment
...
spec:
  initContainers:
  - command:
    - curl demo-mysql
  containers:
  - env:
    - name: WORDPRESS_DB_HOST
      value: demo-mysql
```

Use Case #3 ConfigMaps generated from multiple sources

归并属性

base



production overlay



/dev/stdout

kustomization.yaml

```
configMapGenerator:
- name: myCMap
  files:
  - common.properties
```

common.properties

```
color=blue
height=10m
```

kustomization.yaml

```
bases:
- ../../base
namePrefix: prod-
configMapGenerator:
- name: myCMap
  behavior: merge
  files:
  - secret.properties
```

secret.properties

```
dbpassword=foo
```

```
kind: ConfigMap
metadata:
  name: prod-myCMap-b5m75cxc
data:
  color=blue
  height=10m
  dbpassword=foo
```

... so on for **staging** and **development** variants.

Properties can be owned by *different* teams.

It's all *patching*.

kustomize is just a means to manage
k8s-targeted patching.

Things one might want to customize:

你可能想配置以下内容

context namespaces, names, labels


container image tag, args, env, config files, secrets, static data

budgets replicas, cpu, memory, volume source

policies RBAC, pod security, network

Use Case #4 Deploy from version control 从版本控制进行部署

```
$ kustomize build \  
  github.com/kubernetes-sigs/kustomize/examples/helloWorld |\  
  kubectl apply -f -
```



Use Case #4 Deploy from version control 从版本控制进行部署

This deploys from a git tag,
merely adding a name prefix.

kustomization.yaml

```
...  
  
namePrefix: hello-  
bases:  
- github.com/kubernetes-sigs/kustomize/examples/multibases?ref=v1.0.6
```



Use Case #5 Rollback

回滾

```
$ cd target
$ git checkout HEAD~1
$ kustomize build . | kubectl apply -f -
```

Verify: `kustomize build . | kubectl diff -f -`

And/or use `kustomization.yaml` to add git tag annotations to all resources, then verify rollback with `kubectl describe`.

Use Case #6 Keep up with upgrades

如何和别人的配置升级保持一致？

Fork their config.

Add a `kustomization.yaml` to adapt it to your needs.

Periodically *git rebase* to keep up.

No merge conflicts ever → *upgrade can be automated.*



A git repo of resources,
with **kustomizations** describing **variants**,
is a simple, powerful way to manage
shared configuration **patching**.

yourApp

```
|— README.md
|— base
|   |— kustomization.yaml
|   |— deployment.yaml
|   |— env.txt
|   |— service.yaml
|— overlays
|   |— production
|   |   |— kustomization.yaml
|   |   |— replica_count.yaml
|   |   |— cpu_count.yaml
|   |— staging
|   |   |— kustomization.yaml
|   |   |— ...
|— ...
```

DEMO


No - too boring

不要啊！好无聊！

It's just YAML to stdout! Try these [examples](#)

kustomize design goals
kustomize的设计目标

Talk not over yet...
almost!




composable 可组合性

plain text

do one thing

pipe friendly

say nothing



In 1973 Douglas McIlroy encourages Ken Thompson to add pipes to unix. Doug wrote diff, tee, tr, echo, sort, etc.

```
$ kustomize build $target | kubectl apply -f -
```

k8s style k8s风格

patch raw yaml resources

extensible to CRDs

targets *kubectl apply*

use patch concepts already in *apply*

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 4
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.8
          ports:
            - containerPort: 80
  ...
```


git workflow

工作流

declarative YAML only

diff cluster against git

share bases via git

capture upgrades with rebase

```
# config upgrade script
```

```
$ cd yourConfigDirectory
```

```
$ git fetch upstream
```

```
$ git rebase upstream/master
```

```
$ git push origin master
```

simple
简单

no templating

no new language

no unintentional API

no foreign notion of a
'package' or 'application'

```
// [1] http://google3/production  
all_packages = filter(lambda y
```

```
list_of_package_names_map =  
    flatten(map(lambda x:  
        encode_list
```

```
all_packages  
[binary_packages
```

```
packages_map_as_string =  
    pkg_csum_errors = cond
```

```
// Note: Do not call  
// word.  
packages_list_tmp =  
    flatten([all_packages  
        cond(de
```

```
// Clear the field  
// have a Borg package  
// that have not  
packages_list =
```

```
error string
```

```
kind: Deployment  
metadata:  
  name: {{ template "artifactory.fullname" . }}  
  labels:  
    app: {{ template "artifactory.name" . }}  
    chart: {{ .Chart.Name }}-{{ .Chart.Version }}  
    component: "{{ .Values.artifactory.component }}"  
    heritage: {{ .Release.Service }}  
    release: {{ .Release.Name }}  
spec:  
  replicas: {{ .Values.artifactory.replicas }}  
  template:  
    metadata:  
      labels:  
        app: {{ template "artifactory.name" . }}  
        component: "{{ .Values.artifactory.component }}"  
        release: {{ .Release.Name }}  
    spec:  
      {{- if .Values.imagePullSecrets }}  
      imagePullSecrets:  
        - name: {{ .Values.imagePullSecrets.name }}  
      {{- end }}  
      initContainers:  
        - name: "remove-lost-found"  
          image: "{{ .Values.initContainerImage }}"  
          imagePullPolicy: {{ .Values.artifactory.imagePullPolicy }}  
          command:  
            - 'sh'  
            - '-c'  
            - 'rm -rfv {{ .Values.artifactory.persistentVolumeMounts.mountPath }}'  
      volumeMounts:  
        - mountPath: {{ .Values.artifactory.persistentVolumeMounts.mountPath }}  
          name: artifactory-volume
```

```
1 function create_alpine_pod()  
2   local pod = {  
3     apiVersion = "v1",  
4     kind = "Pod",  
5     metadata = {  
6       name = alpine_fullname(_),  
7       labels = {  
8         heritage = _ .Release.Service or "",  
9         release = _ .Release.Name,  
10        chart = _ .Chart.Name .. "-" .. _ .Chart.Version,  
11        app = alpine_name(_)  
12      }  
13    },  
14    spec = {  
15      restartPolicy = _ .Values.restartPolicy,  
16      containers = {  
17        {  
18          name = waiter,  
19          image = _ .Values.image.repository,  
20          imagePullPolicy = _ .Values.imagePullPolicy,  
21          command = {  
22            "/bin/sleep",  
23            "9000"  
24          }  
25        }  
26      }  
27    }  
28  }  
29  
30  _ .resources.add(pod)  
31 end
```

Already integrated into [skaffold](#), [kubebuilder](#), [Replicated Ship](#), ...

Coming soon:

kubectl integration

\$ kubectl apply -f target

← will honor
kustomization.yaml
if present

example site

like [godoc.org](#), but for kustomizations

Thanks! 非常感谢！

contributors

<- please file issues / help

replicated

kustomize.io!

sig-cli

sponsoring kustomize!

brian grant

principal eng @google, k8s founder,
author of declarative application management in
Kubernetes, the inspiration for kustomize.

Questions 提问时间



backup slides

Template / DSL
paradox:

*The more people you try to please,
the worse it gets.*

Helm Strategy

1. Download *someChart*
2. `$ helm template someChart --output_dir target`
3. Add your own `kustomization.yaml` to `target`.
4. Commit `target` to git.
5. Periodically do (1) again and *manually* merge changes.
...Or not. Do you think it's important?



What would you say you do here?



I find an example.
Then I kustomize it.

operands

resources - file names on disk

generated resources - instructions

CRDs - expands the list of recognized resources

bases - nested kustomizations

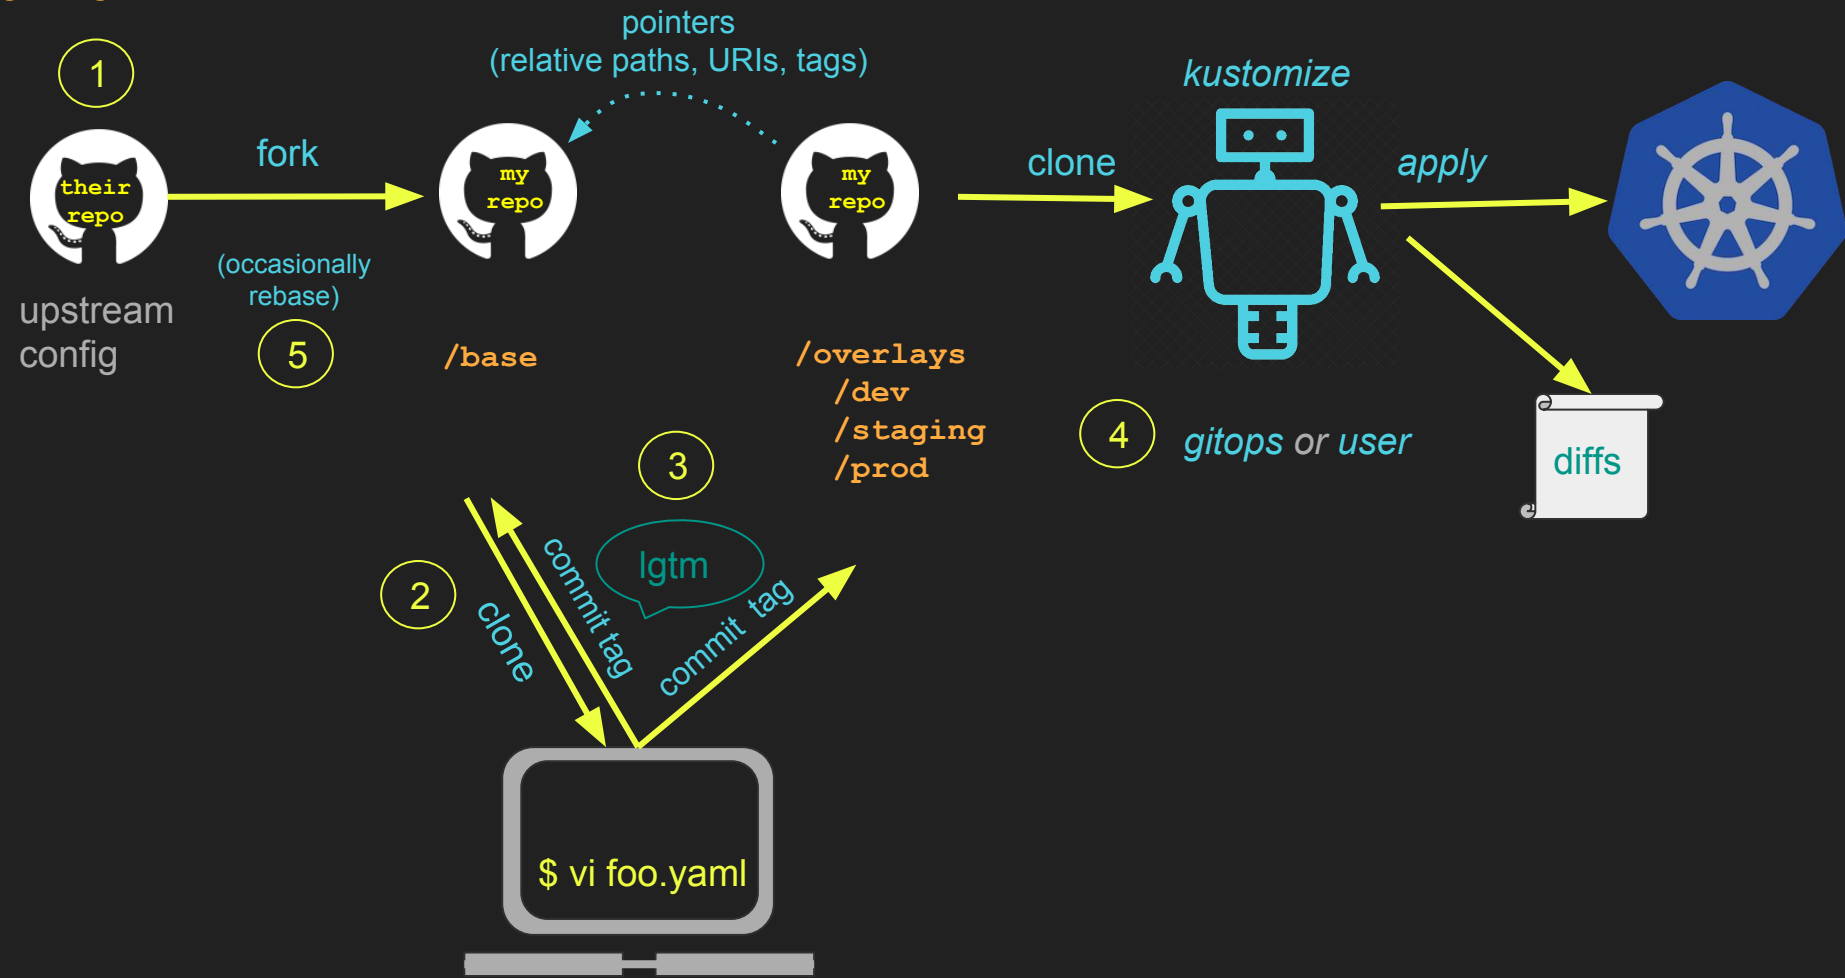
operations

add name prefix

add labels and annotations

patch... (etc.)

workflow



Config sharing:

- Share your resource set with the team / company / world. They kustomize it, and rebase periodically.
- Dev-ops shares common resources to company eng teams. They add their own kustomizations.
- Share common resources to different environments - dev/staging/production, blue/green, etc., kustomizing the differences.